

# Label-Efficient Dark-Vessel Detection in SAR: Does SAR-Domain Pretraining Outperform Optical and ImageNet Transfer Across ViT and CNN?

John Roth      Kyle Wagner

Johns Hopkins University • EN.705.643 Deep Learning Developments with PyTorch • August 2026

## Abstract

Dark vessels operate without a matched Automatic Identification System report. Spaceborne synthetic aperture radar (SAR) can find them through cloud and darkness, but human-verified examples are scarce. We test whether source-domain pretraining reduces the labeled scenes a detector needs on xView3-SAR, and whether the effect persists across architectures. Two size-matched tracks use ViT-B/16 and ConvNeXt-V2-Base backbones. Within each track, four arms compare random initialization with released ImageNet, optical remote-sensing, and SAR checkpoints under one detector, optimizer, schedule, scene-level split, and seed. We fine-tuned nested 10–100% fractions of 111 training scenes, which gives 32 controlled cells. All 32 cells completed on H100 GPUs in strict FP32, and the build re-derives every published value from hash-bound evidence. Transfer helps most when labels are scarce. The SAR-versus-optical contrast depends on the architecture: SAR pretraining beats optical in the CNN track at every budget, but optical leads in the ViT track below half data; the held-out test split confirms this sign at every budget. Development selection favors generic ImageNet at full data in both tracks, but the held-out test split instead favors SAR in the ViT track. Reported values are corrected development-selection F1 alongside once-scored held-out test results. The 50-scene human-verified evaluation remains sealed for follow-on work.

## 1 Introduction

Spaceborne synthetic aperture radar (SAR) provides wide-area maritime observation independent of daylight and through most weather conditions [7]. Those properties make Sentinel-1 imagery useful for detecting vessels without a contemporaneous matched Automatic Identification System report—the so-called *dark vessels*. Such detections support maritime safety, fisheries monitoring, and investigation of activity potentially associated with sanctions evasion. The task is difficult: vessels occupy few pixels in a large scene, sea clutter and coastal structures produce strong responses, and human-verified labels are expensive. The xView3-SAR benchmark makes the problem concrete with ground-range-detected imagery, georeferenced point annotations, and an evaluation protocol based on geographic matching [3].

The scarcity of verified labels motivates transfer learning, but which source helps most is unknown. Generic ImageNet pretraining supplies broad visual features; optical remote-sensing pretraining adds overhead-scene statistics; SAR pretraining also shares the sensing modality. These alternatives interact with architecture. Vision transformers and convolutional networks encode different inductive biases, and no released recipe matches the remote-sensing checkpoints method for method. A comparison that changes architecture, optimization, detector, and pretraining source at the same time cannot isolate the source effect.

We study label efficiency through two architecture-matched tracks (Fig. 3, Appendix B). Within each track, four arms differ only in initialization: random, ImageNet, optical remote sensing, or SAR. Every arm fine-tunes end to end through the same point detector, input representation, optimizer, schedule, nested scene subsets, and fixed seed, producing 32 controlled cells. The study asks three questions:

1. How does random, ImageNet, optical, or SAR initialization change detection F1 as the number of labeled scenes grows?
2. Does SAR pretraining beat optical pretraining within each architecture, especially under label scarcity, and is the direction of that contrast consistent across tracks?
3. After model selection and thresholds freeze, how do the arms generalize to held-out scenes and human-verified dark vessels?

This report answers the first two questions with the completed 32-cell H100 campaign; the build re-derives every value from hash-bound evidence. The once-scored 16-scene test matrix addresses the held-out part of

Table 1: Frozen scene-level data protocol and annotation support under the corrected ground-truth contract. The 50 verified scenes remain sealed until the held-out test matrix and review gates complete.

| Partition                | Scenes | Vessel H/M             | Non-vessel H/M | LOW ignore |
|--------------------------|--------|------------------------|----------------|------------|
| Training-time dev subset | 8      | 517                    | 107            | 118        |
| Development (full)       | 23     | 1,479                  | 804            | 441        |
| Test                     | 16     | 1,165                  | 420            | 325        |
| Verified final           | 50     | sealed; human verified |                |            |

the third. Our contribution is experimental control rather than a new method. We provide an eight-arm two-track design in which initialization is the only within-track variable, one corrected evaluation contract with checkpoint-bound operating thresholds, and a fail-closed evidence path from training bytes to every published number.

## 2 Dataset and Evaluation Protocol

**Study cohort.** We use xView3-SAR Sentinel-1 ground-range-detected scenes with VH and VV polarizations at 10 m pixel spacing [3]. Scene identifiers define every split. A frozen seed-0 cohort contains 150 study scenes: 111 train, 23 development, and 16 test, stratified by coarse geographic region and shoreline presence. The protocol reserves a separate set of 50 human-verified scenes for one sealed final evaluation. The four training budgets are exact nested sets of 12, 28, 56, and 111 scenes:

$$S_{10} \subset S_{25} \subset S_{50} \subset S_{100}. \quad (1)$$

Every arm sees the same scene set at a given budget, so a budget never changes which scenes one initialization sees relative to another. The preprocessing path converts each scene into a fixed three-channel decibel-domain tensor [VH, VV, VH − VV]. It computes normalization statistics from training scenes only. It then cuts overlapping 800-pixel chips, and optimization samples shared 512-pixel crops from those chips.

**Ground-truth contract.** One target-semantics contract serves training and evaluation (Fig. 5, Appendix B). A row is a positive vessel point only when `is_vessel` is true at HIGH or MEDIUM confidence. HIGH/MEDIUM non-vessels are background. LOW-confidence rows define ignore regions. Table 1 lists the annotation support of every partition under this contract; a committed audit receipt binds the counts to the frozen split file by SHA-256. These support counts are pre-registered integrity checks: evaluation stops if regenerated counts differ. Training never stages test or human-verified labels. Appendix B (Fig. 2) shows the scale of one study scene against a single chip and the scene centroids of the frozen splits.

**Selection and scoring.** The frozen scorer greedily matches predictions to ground truth within 200 m and aggregates precision, recall, and F1 across scenes. Every five epochs, each run performs full-scene inference on the same eight development scenes and sweeps the score threshold that maximizes development F1. The run records the best value with the retained checkpoint. The protocol binds that complete operating point (epoch, threshold, and counts) to the SHA-256 of the exact checkpoint. Held-out stages apply each stored threshold verbatim, first to the 16 test scenes and then, once, to the human-verified set. Held-out annotations never retune a threshold or select a checkpoint. The frozen split is leakage-free but not geographically independent: Sentinel-1 revisits place held-out scenes near training scenes, so we describe absolute performance as in-region generalization.

## 3 Related Work

xView3 frames dark-vessel detection as point localization in Sentinel-1 SAR rather than box classification [3]. A CenterNet-style heatmap fits the point annotations and distance metric without inventing box extents [9]. We use this established detector form to hold the task head fixed, and we do not present the detector as a contribution.

ImageNet pretraining is the common baseline for visual transfer learning [4]. The ViT track’s generic control is a supervised AugReg ViT-B/16 checkpoint; the CNN track’s is ConvNeXt-V2-Base initialized by fully convolutional masked autoencoding followed by supervised ImageNet-1K fine-tuning [8]. Remote-sensing pretraining may narrow the gap between source and target statistics: SatDINO applies DINO self-distillation to fMoW-RGB optical imagery [5], and SARMAE applies masked autoencoding to a million-scale

Table 2: Controlled initialization roles. All non-random entries are released, revision-pinned checkpoints; task heads are discarded.

| Role     | ViT-B/16 track            | ConvNeXt-V2-Base track    |
|----------|---------------------------|---------------------------|
| Random   | fresh initialization      | fresh initialization      |
| Optical  | SatDINO, DINO on fMoW-RGB | BigEarthNet-v2 Sentinel-2 |
| SAR      | SARMAE, MAE on SAR-1M     | BigEarthNet-v2 Sentinel-1 |
| ImageNet | supervised AugReg         | FCMAE + supervised        |

SAR corpus [2]. These released ViT checkpoints differ in objective as well as sensing domain, so their contrast compares the best available checkpoints rather than modality alone. The CNN track is tighter. Its optical and SAR arms use BigEarthNet-v2 land-cover models trained with the same architecture, dataset family, and supervised task on Sentinel-2 and Sentinel-1 inputs [6, 1]. That pair is the study’s cleanest modality contrast. Matched roles across the two tracks then test whether a within-track pattern survives a change of architecture family.

## 4 Methods

**Two matched tracks, eight arms.** Table 2 lists the arms. Within each track, the architecture and parameter count stay fixed and only the initialization changes. Pretrained checkpoints transfer the encoder only, with source heads discarded. Random floors use the identical target architecture. End to end, the implemented detectors contain 89,996,801 (ViT) and 93,988,865 (CNN) trainable parameters.

**Shared point detector.** Every arm predicts a stride-4 vessel-presence heatmap from the same three-channel input (Fig. 4, Appendix B). Gaussian targets with standard deviation two output pixels mark positives; LOW-confidence annotations define ignore neighborhoods. Training uses the penalty-reduced focal objective of CenterNet-style point detection [9]. The ViT adapter receives stride-16 features and the CNN adapter stride-32 features with one learned upsampling stage; both terminate in the same 256-channel head. At inference, local peaks above a permissive 0.05 candidate floor map to scene coordinates, and 120 m distance non-maximum suppression removes duplicates. The tiler cuts whole scenes into 512-pixel windows at stride 384.

**Uniform fine-tuning recipe.** One PyTorch Lightning module, data module, and trainer configuration serve all cells. AdamW uses learning rate  $10^{-4}$ , weight decay 0.05, layer-wise decay 0.65, five warm-up epochs, and cosine decay across at most 50 epochs; early stopping waits four development evaluations. Every cell uses micro-batch 16, accumulation 1, seed 0, and full-encoder fine-tuning. The campaign executed one process per NVIDIA H100 in Lightning `32-true` with CUDA matmul and cuDNN TF32 disabled. The code itself is hardware-neutral; the uniform hardware class is a property of the evidence. No arm and no label fraction gets its own hyperparameters.

**Registered contrasts.** For track  $t$ , role  $r$ , and fraction  $f$ , let  $F_{r,f}^{(t)}$  denote development F1 at that cell’s checkpoint-bound operating threshold. The analysis reports transfer over the track-matched random floor,

$$\Delta_{r,f}^{(t)} = F_{r,f}^{(t)} - F_{\text{random},f}^{(t)}, \quad (2)$$

and the within-track SAR–optical contrast,

$$\Delta_{\text{SAR-opt},f}^{(t)} = F_{\text{SAR},f}^{(t)} - F_{\text{optical},f}^{(t)}. \quad (3)$$

We assess architecture generality from agreement in the sign and pattern of Eq. (3) across tracks, and we never pool the tracks into one effect estimate. The evidence generator rejects a marker that does not hash to its cohort binding, a metric inconsistent with its own TP/FP/FN counts, a best-development value that disagrees with the training curve, a mixed hardware class, or a partial held-out cohort.

## 5 Results and Analysis

**Execution summary and reporting contract.** The core result is one complete cohort: eight arms at four nested budgets, seed 0, all 32 cells finished on one NVIDIA H100 80GB HBM3 class in strict FP32

Table 3: Development-selection F1 for all 32 cells at each cell’s checkpoint-bound operating threshold. Budgets are fractions of the 111 training scenes (scene counts in parentheses). Per-cell thresholds, best epochs, and held-out slots appear in Appendix A.

| Track         | Initialization | 10% (12) | 25% (28) | 50% (56) | 100% (111) |
|---------------|----------------|----------|----------|----------|------------|
| ViT-B/16      | Random         | 0.8030   | 0.8563   | 0.8912   | 0.8958     |
|               | Optical        | 0.8902   | 0.8802   | 0.9136   | 0.9268     |
|               | SAR            | 0.8857   | 0.8624   | 0.9161   | 0.9289     |
|               | ImageNet       | 0.8761   | 0.8716   | 0.9284   | 0.9399     |
| ConvNeXt-V2-B | Random         | 0.7992   | 0.8167   | 0.8664   | 0.8919     |
|               | Optical        | 0.7992   | 0.8063   | 0.8586   | 0.8636     |
|               | SAR            | 0.8762   | 0.8868   | 0.8998   | 0.9068     |
|               | ImageNet       | 0.8856   | 0.9005   | 0.9221   | 0.9387     |




Figure 1: Development-F1 label-efficiency curves, one panel per track. Solid lines are pretrained arms; the dashed gray line is the track’s random floor. Values are seed-0 point estimates; no uncertainty ribbon is drawn.

(555.8 GPU-hours; code [1a82d508](#); cohort [9b1ba03e](#)). The build re-derives every published value from the committed evidence: the frozen cohort, the hash-bound completion markers, and the full training curves. A single inconsistent byte stops the build. The reported metric is corrected development-selection F1. Each cell selects its checkpoint and operating threshold on the eight development scenes, so the values carry selection optimism and are not held-out estimates. Table 3 lists the matrix; Fig. 1 shows the four-point curves. Training dynamics appear in Appendix B (Fig. 6).

**Label-efficiency curves and transfer gains.** The two tracks occupy different regimes. Every pretrained ViT arm clears its random floor at 12 scenes: optical by +0.087, SAR by +0.083, and ImageNet by +0.073 (Eq. (2); gain panels in Fig. 7, Appendix B). The CNN track spreads wider. SAR gains +0.077 and ImageNet +0.086 over the 12-scene floor. The optical Sentinel-2 arm only matches its floor at 12 scenes (both 0.7992) and falls below it at 25% and 50%: the land-cover source gives this backbone features that fine-tuning must first undo.

Two curves violate monotonicity at the 10%–25% transition: ViT-optical drops from 0.8902 to 0.8802 and ViT-SAR from 0.8857 to 0.8624. We report these declines as observed, and we do not smooth or rerun them. The metric rests on a few hundred positives over eight fixed scenes, so a shift of a few hundredths corresponds to a handful of detections at the swept threshold. The protocol resolves such selection noise with held-out scoring on the sealed sixteen-scene split.

We also ask for the smallest tested budget at which an arm reaches its track’s full-data random floor. CNN-ImageNet crosses at 28 scenes (0.9005 against 0.8919). Every ViT transfer arm and CNN-SAR cross at 56 scenes. CNN-optical never crosses. No arm crosses at 12 scenes: the strongest 12-scene cell, ViT-optical at 0.8902, sits just below the ViT floor’s 0.8958. Under this development-selection metric, a well-chosen

public checkpoint at 12 scenes nearly, but not fully, substitutes for 111 random-initialization scenes.

**Source-domain contrasts and architecture generality.** The CNN track orders SAR above optical at every budget (Eq. (3); right panel of Fig. 7), with the largest contrast under scarcity (0.8762 versus 0.7992 at 12 scenes). The BigEarthNet Sentinel-1 and Sentinel-2 checkpoints share architecture, dataset family, and supervised task, so this is the study’s cleanest modality effect. The ViT track shows the opposite order below half data: optical leads SAR at 10% and 25% (0.8902 versus 0.8857; 0.8802 versus 0.8624), and SAR moves ahead only from 56 scenes upward. The sign of the SAR–optical contrast does not replicate across architectures under label scarcity. A modality-matched encoder is not always the best scarce-label choice. Architecture and pretraining objective moderate the effect, and the released ViT checkpoints also differ in objective and source volume, so their contrast describes available checkpoints rather than modality alone.

At full data, generic ImageNet initialization leads both tracks (0.9399 and 0.9387). Domain pretraining matters most where labels are scarce. With enough target data, generic visual structure plus full fine-tuning closes the domain gap.

**Held-out results.** All 32 cells were scored once on the frozen 16-scene test split with their cohort-bound development thresholds (inference precision 32-true), and the generator admits the test column only because all 32 results validated against the cohort. Every cell’s test F1 sits below its development F1; the shrinkage ranges from 0.041 to 0.142 F1 and averages 0.080, consistent with development-selection optimism (both checkpoint and threshold are chosen on eight scenes) rather than a threshold-transfer failure. The SAR–optical sign holds at every budget in both tracks, including the reversed ViT order below half data. One ranking changes: development selection favors ImageNet at full data in both tracks, but held-out scoring instead favors SAR in the ViT track (0.8748 versus ImageNet’s 0.8569); the CNN track’s ImageNet lead survives (0.8966).

The campaign’s grid-validation gate stopped on two non-monotonic test-F1 curves, both in the SAR role. ViT-SAR repeats the dev-side drop reported above: 0.8041 at 10% falls to 0.7715 at 25%. CNN-SAR shows a violation development selection never revealed: its development curve rises from 50% to 100% (0.8998 to 0.9068), but its test curve falls (0.8521 to 0.8221), the study’s clearest case of development selection missing a real instability. Appendix A lists the full test column. The 50 human-verified scenes remain sealed. The protocol defines dark-vessel recall and near-shore performance only on that set, which awaits its once-only evaluation.

## 6 Conclusion and Future Work

We built a controlled two-track experiment for measuring transfer under four nested label budgets: one detector, one corrected label contract, one strict-FP32 recipe, and a fail-closed path from training bytes to every published number. The completed 32-cell grid supports three conclusions. First, transfer concentrates its value under scarcity. At 12 scenes, every pretrained arm except CNN-optical clears its random floor, and the best 12-scene cells approach the full-data random floors. Second, matching the sensing modality is not always the right scarce-label choice. SAR pretraining beats optical in the CNN track at every budget, but optical leads in the ViT track below half data, so the contrast’s sign depends on the architecture (the held-out test split confirms this sign at every budget). Third, generic ImageNet initialization wins both tracks in development selection at full data, though held-out scoring instead favors SAR in the ViT track. These remain corrected development-selection results from one seed and support no variance claim.

The once-scored 16-scene test column (mean change 0.080 F1; Sect. 5) confirms the scarcity and SAR-versus-optical conclusions but reverses the full-data ImageNet-versus-SAR ranking in the ViT track: development selection alone would have missed that reversal. The 50 human-verified scenes then support one sealed evaluation of dark-vessel recall and near-shore performance for follow-on work.

**Roles.** John Roth and Kyle Wagner jointly designed the study, implemented its data, model, training, and evaluation code, ran the experiments, analyzed the evidence, and wrote this report.

## References

- [1] K. N. Clasen, L. Hackel, T. Burgert, G. Sumbul, B. Demir, and V. Markl. reBEN: Refined BigEarthNet dataset for remote sensing image analysis. In *Proceedings of the IEEE International Geoscience and Remote Sensing Symposium*, pages 1264–1268, 2025.
- [2] D. Liu, D. Wang, H. Wang, et al. SARMAE: Masked autoencoder for SAR representation learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 6496–6507, 2026.
- [3] F. S. Paolo, T.-t. T. Lin, R. Gupta, B. Goodman, N. Patel, D. Kuster, D. Kroodsmas, and J. Dunnmon. xView3-SAR: Detecting dark fishing activity using synthetic aperture radar imagery. In *Advances in Neural Information Processing Systems*, volume 35, pages 37604–37616, 2022.
- [4] A. Steiner, A. Kolesnikov, X. Zhai, R. Wightman, J. Uszkoreit, and L. Beyer. How to train your ViT? data, augmentation, and regularization in vision transformers. *Transactions on Machine Learning Research*, 2022.
- [5] J. Straka and I. Gruber. SatDINO: A deep dive into self-supervised pretraining for remote sensing, 2025.
- [6] G. Sumbul, M. Charfuelan, B. Demir, and V. Markl. BigEarthNet: A large-scale benchmark archive for remote sensing image understanding. In *Proceedings of the IEEE International Geoscience and Remote Sensing Symposium*, pages 5901–5904, 2019.
- [7] R. Torres, P. Snoeij, D. Geudtner, et al. GMES Sentinel-1 mission. *Remote Sensing of Environment*, 120:9–24, 2012.
- [8] S. Woo, S. Debnath, R. Hu, X. Chen, Z. Liu, I. S. Kweon, and S. Xie. ConvNeXt V2: Co-designing and scaling ConvNets with masked autoencoders. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 16133–16142, 2023.
- [9] X. Zhou, D. Wang, and P. Krähenbühl. Objects as points. *arXiv preprint arXiv:1904.07850*, 2019.

## A Full Evidence Inventory

Table 4 lists every cell: development-selection F1, the checkpoint-bound operating threshold, the epoch of the selected checkpoint, the epochs run, measured GPU-hours, and held-out test F1 where admitted. The build re-derives all values from the committed cohort, markers, training curves, and runtime receipts. Dashes mean the fail-closed generator admitted no value, never zero performance.

Table 4: Per-cell evidence inventory (development-selection values; the test column is admitted only when all 32 immutable held-out results validate).

| Cell         | Budget | Dev F1 | Threshold | Best ep. | Epochs | GPU-h | Test F1 |
|--------------|--------|--------|-----------|----------|--------|-------|---------|
| ViT random   | 10%    | 0.8030 | 0.822     | 19       | 40     | 3.5   | 0.7042  |
| ViT random   | 25%    | 0.8563 | 0.890     | 29       | 50     | 8.1   | 0.7372  |
| ViT random   | 50%    | 0.8912 | 0.853     | 44       | 50     | 14.8  | 0.8101  |
| ViT random   | 100%   | 0.8958 | 0.778     | 14       | 35     | 19.8  | 0.8101  |
| ViT optical  | 10%    | 0.8902 | 0.888     | 49       | 50     | 4.3   | 0.8198  |
| ViT optical  | 25%    | 0.8802 | 0.910     | 24       | 45     | 7.3   | 0.8279  |
| ViT optical  | 50%    | 0.9136 | 0.902     | 39       | 50     | 14.8  | 0.8495  |
| ViT optical  | 100%   | 0.9268 | 0.859     | 34       | 50     | 28.0  | 0.8465  |
| ViT SAR      | 10%    | 0.8857 | 0.587     | 4        | 25     | 2.2   | 0.8041  |
| ViT SAR      | 25%    | 0.8624 | 0.768     | 4        | 25     | 4.1   | 0.7715  |
| ViT SAR      | 50%    | 0.9161 | 0.855     | 39       | 50     | 14.8  | 0.8672  |
| ViT SAR      | 100%   | 0.9289 | 0.904     | 24       | 45     | 25.2  | 0.8748  |
| ViT ImageNet | 10%    | 0.8761 | 0.906     | 39       | 50     | 4.2   | 0.8103  |
| ViT ImageNet | 25%    | 0.8716 | 0.929     | 34       | 50     | 8.1   | 0.8051  |
| ViT ImageNet | 50%    | 0.9284 | 0.844     | 49       | 50     | 14.8  | 0.8545  |
| ViT ImageNet | 100%   | 0.9399 | 0.894     | 49       | 50     | 28.1  | 0.8569  |
| CNN random   | 10%    | 0.7992 | 0.716     | 29       | 50     | 7.0   | 0.6571  |
| CNN random   | 25%    | 0.8167 | 0.819     | 44       | 50     | 14.2  | 0.7065  |
| CNN random   | 50%    | 0.8664 | 0.857     | 44       | 50     | 26.6  | 0.7807  |
| CNN random   | 100%   | 0.8919 | 0.910     | 49       | 50     | 51.4  | 0.8124  |
| CNN optical  | 10%    | 0.7992 | 0.749     | 29       | 50     | 6.8   | 0.6712  |
| CNN optical  | 25%    | 0.8063 | 0.822     | 44       | 50     | 14.1  | 0.7024  |
| CNN optical  | 50%    | 0.8586 | 0.879     | 34       | 50     | 26.5  | 0.7616  |
| CNN optical  | 100%   | 0.8636 | 0.878     | 44       | 50     | 51.3  | 0.7948  |
| CNN SAR      | 10%    | 0.8762 | 0.776     | 34       | 50     | 6.7   | 0.8000  |
| CNN SAR      | 25%    | 0.8868 | 0.877     | 14       | 35     | 9.8   | 0.8025  |
| CNN SAR      | 50%    | 0.8998 | 0.893     | 24       | 45     | 23.9  | 0.8521  |
| CNN SAR      | 100%   | 0.9068 | 0.832     | 4        | 25     | 25.8  | 0.8221  |
| CNN ImageNet | 10%    | 0.8856 | 0.869     | 4        | 25     | 3.4   | 0.7810  |
| CNN ImageNet | 25%    | 0.9005 | 0.906     | 39       | 50     | 14.0  | 0.8481  |
| CNN ImageNet | 50%    | 0.9221 | 0.899     | 19       | 40     | 21.2  | 0.8807  |
| CNN ImageNet | 100%   | 0.9387 | 0.894     | 44       | 50     | 51.2  | 0.8966  |

The campaign’s sanitized operator status snapshot (captured before the final cells closed) remains committed beside the evidence tree as a historical record of execution state at the deadline.

## B Study Data and Figures




Figure 2: Study data. Top: scale context for one study scene, a full Sentinel-1 ground-range-detected scene and the single  $800 \times 800$  chip ( $8 \times 8$  km) marked by the red box; vessels occupy only a few pixels at this scale. Bottom: scene centroids of the frozen splits, computed from the label records. Sentinel-1 revisits place held-out scenes near training scenes, which is why Sect. 2 reports in-region generalization.




Figure 3: Controlled study design. Initialization is the only variable within each architecture track. Matched roles are compared descriptively across tracks; the final human-verified set remains sealed until the cohort and analysis code are frozen.






Figure 4: The shared detection path: only the encoder differs by track; one head, decode rule, dev-selected threshold, and frozen scorer serve every arm.




Figure 5: The corrected ground-truth contract on the training-time development subset. A row is a positive only when `is_vessel` is true at HIGH or MEDIUM confidence; HIGH/MEDIUM non-vessels are background, and LOW-confidence rows become ignore regions. Every H100 cell trained and evaluated under this contract.




Figure 6: Development-F1 evaluation history for the smallest (12 scenes, solid) and largest (111 scenes, dotted) budgets, with each cell’s selected checkpoint marked. Pretrained CNN arms separate from the random floor within the first evaluations; several strong arms peak early, and the checkpoint binding retains those early peaks.




Figure 7: Left, center: transfer gain over the track-matched random floor (Eq. (2)) at each budget; the optical CNN arm is at or below zero everywhere. Right: the SAR-optical contrast (Eq. (3)); its sign flips between tracks below half data, so the modality effect is architecture-dependent.




Figure 8: Operating points for all 32 cells (filled circles: ViT; open squares: CNN; colors identify roles as in Fig. 1). Left: precision-recall positions with iso-F1 contours. Center: swept threshold by initialization role. Right: development F1 against measured GPU-hours per cell.

## C Extended Analysis

**Operating thresholds and calibration.** Each cell selects its operating point on the training-time development subset, and the cohort binds that point to the selected checkpoint. Figure 8 shows all 32 of them. The recorded thresholds vary widely (center panel): from 0.587 to 0.929 across the ViT cells and from 0.716 to 0.910 across the CNN cells. The spread confirms that a single shared threshold would misrepresent most arms and that the per-cell sweep does real calibration work. At the bound operating point, precision exceeds recall in 15 of the 32 cells (left panel). The protocol also defines a calibration-transfer diagnostic (the bound threshold against an oracle test-optimal threshold); it becomes available only when the test matrix runs.

**Training dynamics and cost.** Figure 6 traces the development-F1 evaluation history for the smallest and largest budgets. The pretrained CNN arms open a large gap in the first evaluations and hold it, while the CNN random floor climbs for the full schedule; the ViT arms bunch together early. Early stopping (patience of four development evaluations) ended 11 of the 32 cells before the 50-epoch cap. The median best-checkpoint epoch is 34 for the ViT track and 34 for the CNN track. Fine-tuning cost differs between the families: a core cell averaged 12.6 GPU-hours in the ViT track and 22.1 GPU-hours in the CNN track, with the full matrix consuming 555.8 GPU-hours. Within each track, cost scales with the label fraction, so the scarce-label cells that carry the study’s headline contrasts are also its cheapest (right panel of Fig. 8).




Figure 9: A vessel-and-wake target and the shared detector’s heatmap response, produced during development-phase probe runs on this study’s frozen splits (peak score 0.22). Illustrative only; not a scored sample.

**Qualitative example.** Figure 9 shows a representative detector response on a development scene: a heatmap peak localized on a clear vessel-and-wake target at a raw score of 0.22. Peaks this far below 1.0 on clear targets explain two protocol choices, the permissive 0.05 candidate floor and the per-cell threshold sweep; a fixed high threshold would silently drop true detections. The panel illustrates a mechanism, not a scored sample. The registered false-positive taxonomy review is reserved for the held-out stage, and we interpret the panel alongside, not above, the quantitative findings in Sect. 5.

## D Reproducibility and Audit Record

The evidence tree records the frozen training cohort, each cell’s hash-bound completion marker, its full training curve, a sanitized runtime receipt (H100 hardware class, source SHA, strict-FP32 backend state), and the ground-truth audit receipt that binds every support count in Table 1 to the frozen split file. The generator validates marker-to-cohort hashes, metric consistency, and curve agreement before it writes a manuscript artifact. Checkpoint bytes stay outside the repository; we publish their SHA-256 bindings for archive verification. The submission excludes imagery, labels, weights, checkpoints, run trees, credentials, and private paths.

The detector uses 512-pixel crops, a stride-4 heatmap, Gaussian standard deviation two output pixels, a candidate floor of 0.05, and 120m distance NMS. Training uses at most 50 epochs with development

evaluation every five epochs. The project fixes seed 0 and reports point estimates without uncertainty bands.